



# HAProxy Ingress Controller for Kubernetes

by Andjelko Iharos | Dec 12, 2017 | KUBERNETES, TECH | 3 comments

Cloud-based applications have seen a great uptake in recent years, and that is especially true for microservices-based apps and related orchestration frameworks.

These types of applications create needs for load balancing in new contexts, and here at HAProxy Technologies, we are always happy to empower our community with solutions to help it grow. We have developed a number of features over the past year that enable easier integration of HAProxy with Kubernetes and other dynamic environments. The relevant features include [hitless reloads](#), dynamic configuration without reloading using the [HAProxy Runtime API](#), and [DNS records for service discovery](#).

With HAProxy already being well known for its reliability, extensible features, advanced security, and leadership in performance amongst all free and commercial software load balancers, the Kubernetes user community is also supplementing our efforts with several Ingress Controller implementations that use HAProxy at their core, and we are committed to providing them with expert advice and code contributions to help them make the best use of HAProxy's most advanced features.

In this blog post we are going to show you how HAProxy can improve performance and out-of-the-box features for microservices-based applications by using one of these HAProxy Ingress Controllers.

## Getting Traffic into Kubernetes


**HAProxy**
[COMPANY](#)
[CONTACT US](#)
[GET HAProxy](#)
[GET HAProxy](#)

[BLOG](#)
[CUSTOMER LOGIN](#)

Kubernetes started out by relying on LBaaS capabilities of cloud container engines like GKE and AWS to provide connectivity between users and microservices components. But the community soon realized that there was a need for a more dynamic, independent, and portable way of providing the same functionality. Hence, the Kubernetes Ingress project was started.

Ingress provides an implementation-independent definition of rules that govern the flow of traffic between users and components of a service. This feature is most easily demonstrated through L7 routing. With mappings of end user consumable URLs to specific microservices, Ingress gives users the choice of Ingress Controller that will actually implement the routing. While the option to use GKE or AWS always exists, standalone Kubernetes deployments have gained the option to use any load balancer with an accompanying Ingress Controller.

[PRODUCTS](#)
[SOLUTIONS](#)
[SUPPORT](#)
[RESOURCES](#)

## Dynamic Scaling with HAProxy

The Ingress Controller provides functionality required to satisfy Ingress rules, and it also has to contend with the dynamic nature of the microservices for which it is routing the traffic.

HAProxy is extremely fast and resource-efficient allowing you to get the most out of your infrastructure and minimize latencies in high-traffic scenarios. It also brings an almost endless list of options for tuning and customization. HAProxy's features like dynamic scaling and reconfiguration without reloading are also very valuable in this use case as Kubernetes pods are often spawned, terminated, and migrated in quick bursts and in high amounts, especially during deployments.

In our previous blog posts titled "[Dynamic Scaling for Microservices with the HAProxy Runtime API](#)" and "[DNS for Service Discovery in HAProxy](#)" we have covered two possible methods for dynamic scaling. In this blog post we are going to take a look at how these play a role in a Kubernetes HAProxy Ingress Controller implementation.

We will use the HAProxy Ingress Controller implementation available at [jcmoraisjr/haproxy-ingress](https://github.com/jcmoraisjr/haproxy-ingress). It is a project to which HAProxy Technologies has contributed code that enables the Ingress Controller to take advantage of the HAProxy Runtime API. (Another useful HAProxy Ingress Controller implementation that you could look into would be [appscode/voyager](https://github.com/appscode/voyager).)

## Runtime API in Kubernetes Controller

Starting off with one of the examples provided in the Ingress Controllers documentation, we will deploy and enable two HTTP services with one of them serving as a default fallback.

**HAPROXY**

COMPANY ▾

CONTACT US

GET HAPROXY ▾

GET HAPROXY



BLOG

CUSTOMER LOGIN

This could be done by preparing files `ingress-default-deployment.yaml` and `http-svc-deployment.yaml` as follows:

**ingress-default-deployment.yaml:**

PRODUCTS ▾

SOLUTIONS ▾

SUPPORT ▾

RESOURCES ▾

```
1.  apiVersion: extensions/v1beta1
2.  kind: Deployment
3.  metadata:
4.    labels:
5.      run: ingress-default-backend
6.      name: ingress-default-backend
7.  spec:
8.    replicas: 1
9.    selector:
10.     matchLabels:
11.       run: ingress-default-backend
12.   template:
13.     metadata:
14.       labels:
15.         run: ingress-default-backend
16.     spec:
17.       containers:
18.         - name: ingress-default-backend
19.           image: gcr.io/google_containers/defaultbackend:1.0
20.           ports:
21.             - containerPort: 8080
22.
23. ---
24. apiVersion: v1
25. kind: Service
26. metadata:
27.   labels:
28.     run: ingress-default-backend
29.     name: ingress-default-backend
30.     namespace: default
31. spec:
32.   ports:
33.     - name: port-1
34.       port: 8080
35.       protocol: TCP
36.       targetPort: 8080
```

**HAPROXY**

COMPANY ▾

CONTACT US

GET HAPROXY ▾

GET HAPROXY



BLOG

CUSTOMER LOGIN

PRODUCTS ▾

SOLUTIONS ▾

SUPPORT ▾

RESOURCES ▾

**http-svc-deployment.yaml:**

```
1.  apiVersion: extensions/v1beta1
2.  kind: Deployment
3.  metadata:
4.    labels:
5.      run: http-svc
6.      name: http-svc
7.  spec:
8.    replicas: 2
9.    selector:
10.     matchLabels:
11.       run: http-svc
12.  template:
13.    metadata:
14.      labels:
15.        run: http-svc
16.    spec:
17.      containers:
18.        - name: http-svc
19.          image: gcr.io/google_containers/echoserver:1.3
20.          ports:
21.            - containerPort: 8080
22.
23. ---
24. apiVersion: v1
25. kind: Service
26. metadata:
27.   labels:
28.     run: http-svc
29.   name: http-svc
30.   namespace: default
31. spec:
32.   ports:
33.     - name: port-1
34.       port: 8080
35.       protocol: TCP
```



```
36.         targetPort: 8080
37.     selector:
38.     run: http-svc
```

[COMPANY](#) ▾[CONTACT US](#)[GET HAPROXY](#) ▾[GET HAPROXY](#)[BLOG](#)[CUSTOMER LOGIN](#)[PRODUCTS](#) ▾[SOLUTIONS](#) ▾[SUPPORT](#) ▾[RESOURCES](#) ▾

And applying the configuration:

```
1.  $ kubectl apply -f ingress-default-deployment.yaml
2.  deployment "ingress-default-backend" created
3.  service "ingress-default-backend" created
4.
5.  $ kubectl apply -f http-svc-deployment.yaml
6.  deployment "http-svc" created
7.  service "http-svc" created
```

As a next step, we will set up the Ingress rules for L7 routing with the URL `/app` for hostnames, hostname `foo.bar` routed to the pods running the app `http-svc`, and all other URLs routed to pods running the app `ingress-default-backend`.

This could be done by preparing file `http-svc-ingress.yaml` as follows:

**http-svc-ingress.yaml:**

```
1.  apiVersion: extensions/v1beta1
2.  kind: Ingress
3.  metadata:
4.    name: app
5.  spec:
6.    rules:
7.    - host: foo.bar
8.      http:
9.        paths:
10.     - path: /app
```

**HAPROXY**

```
11.      backend:
12.          serviceName: http-svc
13.          servicePort: 8080
14.      - path: /COMPANY
15.      backend:
16.          serviceName: ingress-default-backend
17.          servicePort: 8080
```

[CONTACT US](#)[GET HAPROXY](#) ▾[GET HAPROXY](#)[BLOG](#)[CUSTOMER LOGIN](#)[PRODUCTS](#) ▾[SOLUTIONS](#) ▾[SUPPORT](#) ▾[RESOURCES](#) ▾

And applying the configuration:

```
1. $ kubectl apply -f http-svc-ingress.yaml
2. ingress "app" created
```

Before we start the HAProxy Ingress Controller service, we need to tune its configuration slightly to enable the option “dynamic-scaling” and to set “backend-server-slots-increment” to “4” in this example:

**haproxy-configmap.yaml:**

```
1.  apiVersion: v1
2.  data:
3.      dynamic-scaling: "true"
4.      backend-server-slots-increment: "4"
5.  kind: ConfigMap
6.  metadata:
7.      name: haproxy-configmap
```

And applying the configuration:

**HAPROXY**

COMPANY ▾

CONTACT US

GET HAPROXY ▾

GET HAPROXY



BLOG

CUSTOMER LOGIN

PRODUCTS ▾

SOLUTIONS ▾

SUPPORT ▾

RESOURCES ▾

The option “backend-server-slots-increment” will allow us to tune Ingress Controller behavior in response to fluctuations in the number of pods running. With the setting of “4” in this example, there will always be up to 4 available slots in a backend definition to accommodate adding extra pods before HAProxy will increase the number of available slots. As pods are shut down and more than 4 slots in a backend definition become available, they will be removed to reduce memory consumption. For larger deployments, this setting should be increased accordingly and setting it to hundreds will have no negative impact.

Finally, we can create the HAProxy Ingress Controller deployment:

### haproxy-ingress-deployment.yaml

```
1.  apiVersion: extensions/v1beta1
2.  kind: Deployment
3.  metadata:
4.    labels:
5.      run: haproxy-ingress
6.      name: haproxy-ingress
7.  spec:
8.    replicas: 1
9.    selector:
10.     matchLabels:
11.       run: haproxy-ingress
12.  template:
13.    metadata:
14.      labels:
15.        run: haproxy-ingress
16.    spec:
17.      containers:
18.      - name: haproxy-ingress
19.        image: quay.io/jcmoraisjr/haproxy-ingress
20.        args:
21.        - --default-backend-service=default/ingress-default-backend
```



```

22.      - --default-ssl-certificate=default/tls-secret
23.      - --configmap=$(POD_NAMESPACE)/haproxy-configmap
24.      - --reload-strategy=native
25.  ports:
26.      - name: http
27.        containerPort: 80
28.      - name: https
29.        containerPort: 443
30.      - name: stat
31.        containerPort: 1936
32.  env:
33.      - name: POD_NAME
34.        valueFrom:
35.          fieldRef:
36.            fieldPath: metadata.name
37.      - name: POD_NAMESPACE
38.        valueFrom:
39.          fieldRef:
40.            fieldPath: metadata.namespace

```

COMPANY ▾

CONTACT US

GET HAPROXY ▾

GET HAPROXY



BLOG

CUSTOMER LOGIN

PRODUCTS ▾

SOLUTIONS ▾

SUPPORT ▾

RESOURCES ▾

And apply the configuration:

```

1. $ kubectl apply -f haproxy-ingress-deployment.yaml
2. deployment "haproxy-ingress" created

```

For this example we will also use an explicit service definition to allow client traffic to reach our Ingress Controller. Alternatively, this functionality can be left to external load balancer setups with access to both public network and internal Kubernetes network (a good place for HAProxy or a LBaaS) while retaining the L7 routing of the Ingress Controller.

**haproxy-ingress-svc.yaml:**



**HAPROXY**

COMPANY ▾

CONTACT US

GET HAPROXY ▾

GET HAPROXY



BLOG

CUSTOMER LOGIN

PRODUCTS ▾

SOLUTIONS ▾

SUPPORT ▾

RESOURCES ▾

```
1.  apiVersion: v1
2.  kind: Service
3.  metadata:
4.    labels:
5.      run: haproxy-ingress
6.    name: haproxy-ingress
7.    namespace: default
8.  spec:
9.    externalIPs:
10.     - 10.245.1.4
11.    ports:
12.     - name: port-1
13.       port: 80
14.       protocol: TCP
15.       targetPort: 80
16.     - name: port-2
17.       port: 443
18.       protocol: TCP
19.       targetPort: 443
20.     - name: port-3
21.       port: 1936
22.       protocol: TCP
23.       targetPort: 1936
24.    selector:
25.      run: haproxy-ingress
```

And apply the configuration:

```
1.  $ kubectl apply -f haproxy-ingress-svc.yaml
```

Checking that the L7 routing works as expected, we can see that the URLs below /app are being served by the echo pod, and all others are served by the default pod.



# HAPROXY

COMPANY ▾

CONTACT US

GET HAPROXY ▾

GET HAPROXY



BLOG

CUSTOMER LOGIN

PRODUCTS ▾

SOLUTIONS ▾

SUPPORT ▾

RESOURCES ▾

```

1.  $ curl -s -XGET -H 'Host: foo.bar' 'http://10.245.1.4:80/app'
2.  CLIENT VALUES:
3.  client_address=10.246.97.6
4.  command=GET
5.  real path=/app
6.  query=nil
7.  request_version=1.1
8.  request_uri=http://foo.bar:8080/app
9.
10. SERVER VALUES:
11. server_version=nginx: 1.9.11 - lua: 10001
12.
13. HEADERS RECEIVED:
14. accept=/*/*
15. host=foo.bar
16. user-agent=curl/7.47.0
17. x-forwarded-for=10.246.97.1
18. BODY:
19. -no body in request-
20.
21. $ curl -s -XGET -H 'Host: foo.bar' 'http://10.245.1.4:80/xyz'
22. default backend - 404

```

Most of the above is entirely familiar to regular users of Kubernetes and Ingress, with the notable exception of adding the two options related to dynamic scaling in the HAProxy configmap.

Taking a look at the HAProxy status page on `external_ip:1936`, we can see the two active pods of `http-svc` and two empty slots.

With this, we can now add and remove pods of a service without HAProxy needing to be reloaded as long as the number of pods is within a multiple of the configured “backend-server-slots-increment” value (“4” in our example). When the set of active pods changes, the Ingress Controller issues commands over the HAProxy Runtime API to update the backend definition. At the same time, it records the new configuration in the configuration file for easier inspection and as a failsafe in the case of an unexpected reload.



# HAProxy

COMPANY ▾

CONTACT US

GET HAProxy ▾

GET HAProxy



BLOG

CUSTOMER LOGIN

PRODUCTS ▾

SOLUTIONS ▾

SUPPORT ▾

RESOURCES ▾

```

1. $ kubectl scale deployment http-svc --replicas=3
2. deployment "http-svc" scaled
3.
4. $ kubectl exec haproxy-ingress-1373648123-dqkkc pidof haproxy
5. 38
6.
7. $ kubectl scale deployment http-svc --replicas=1
8. deployment "http-svc" scaled
9.
10. $ kubectl exec haproxy-ingress-1373648123-dqkkc pidof haproxy
11. 38
12.
13. $ kubectl scale deployment http-svc --replicas=4
14. deployment "http-svc" scaled
15.
16. $ kubectl exec haproxy-ingress-1373648123-dqkkc pidof haproxy
17. 38

```

As mentioned, scaling above or below the initial multiple of “backend-server-slots-increment” will cause HAProxy to adjust the number of free slots for bursting pods.

```

1. $ kubectl scale deploy http-svc --replicas=9
2. deployment "http-svc" scaled
3.
4. $ kubectl exec haproxy-ingress-1373648123-dqkkc pidof haproxy
5. 68

```

## Rolling Updates

DevOps teams are often in charge of releasing new application code in production. To help with that, Kubernetes comes with a tool called “rolling update” which will move the pods to the new desired version.

In essence, it works by starting a new pod with the new version of the application, and once it is up and running, the pod with the old version is shut down. Repeating this operation for each pod that is delivering the application could be tedious, but thanks to its controller, HAProxy can perform rolling updates with no effort.

**HAPROXY**

COMPANY CONTACT US

GET HAPROXY

GET HAPROXY



BLOG

CUSTOMER LOGIN

The command to trigger a rolling update could be as follows (assuming that the version "1.4" of echoserver exists):

PRODUCTS ▾

SOLUTIONS ▾

SUPPORT ▾

RESOURCES ▾

```
1. $ kubectl set image deployment http-svc http-svc=gcr.io/google_containers/echoserver:1.4
```

## Conclusion

HAProxy <3 Microservices!

The recent release of HAProxy version 1.8 contains additional functionality added specifically for using DNS SRV records to implement [dynamic scaling using DNS](#).

While using DNS SRV records for configuring backend servers is not Kubernetes-specific, Kube-DNS is perfectly capable of providing this information to the controller, and our solution for making use of it with the same Ingress Controller is coming soon!

If you would like to use HAProxy Enterprise Edition with Kubernetes and get extra benefits such as a fully supported HAProxy installation, Real Time Dashboard, and management and security-focused enterprise add-ons, please see our [HAProxy Enterprise Edition – Trial Version](#) or [contact us](#) for expert advice.

Stay tuned!

SHARE THIS ARTICLE



Tags: [configuration](#), [Ingress Controller](#), [Kubernetes](#), [microservices](#)



## 3 Comments

[COMMUNITY](#)
[CONTACT US](#)
[GET HAPROXY](#)
[GET HAPROXY](#)

[BLOG](#)
[CUSTOMER LOGIN](#)


**Giovanni Silva** on May 21, 2018 at 4:43 pm

One of the HAproxy ingress community projects with good maturity is

<https://github.com/appscode/voyager>

Is backed by a company and opensource

[Reply](#)


**pedro navajas** on July 8, 2018 at 5:19 pm

can you help me ? I make this Ingress, but the Pod is CrashLoopBackOff. This is the pod's log:

It seems the cluster it is running with Authorization enabled (like RBAC) and there is no permissions for the ingress controller. Please check the configuration

[Reply](#)


**Baptiste Assmann** on November 6, 2018 at 5:35 am

You have to configure RBAC for the ingress controller.

Check this link [https://github.com/jcmoraisjr/haproxy-](https://github.com/jcmoraisjr/haproxy-ingress/tree/master/examples/rbac)

[ingress/tree/master/examples/rbac](https://github.com/jcmoraisjr/haproxy-ingress/tree/master/examples/rbac)

[Reply](#)

This site uses Akismet to reduce spam. [Learn how your comment data is processed.](#)

Comments will display after being approved by the moderator.

## RELATED STORIES



**[On-Demand Webinar] Introduction to HAProxy ACLs: Building Rules for Dynamically Routing Requests, Redirecting Users and Blocking Malicious Traffic**



## The Four Essential Sections of an HAProxy Configuration

[COMPANY](#) ▾[CONTACT US](#)[GET HAPROXY](#) ▾[GET HAPROXY](#)[BLOG](#)[CUSTOMER LOGIN](#)

## Introduction to HAProxy ACLs

[PRODUCTS](#) ▾[SOLUTIONS](#) ▾[SUPPORT](#) ▾[RESOURCES](#) ▾

# HAPROXY

---

+1 (844) 222-4340 [contact@haproxy.com](mailto:contact@haproxy.com)

## Connect With Us

## Solutions

[Load Balancing](#)[High Availability](#)[Application Acceleration](#)[Security](#)[Microservices](#)

## Products

[HAProxy Enterprise Edition](#)[Community vs Enterprise Edition](#)[ALOHA Hardware Appliance](#)[ALOHA Virtual Appliance](#)[Software vs Appliance](#)



CUSTOMER LOGIN

## Company

## SOLUTION

### About Us

## Contact Us

## User Reference

## PRODUCTS

**About Us**

RESOURCES ▾

## Blog

## Sitemap