

# Multi Node Kubernetes Cluster with Vagrant, VirtualBox and Kubeadm



Raj Rajaratnam [Follow](#)

Sep 6, 2018 · 4 min read

TL;DR, Go grab my Vagrantfile at <https://github.com/ecom-integration-ballerina/kubernetes-cluster> and run `vagrant up` to bring a 3 node Kubernetes cluster up and running in VirtualBox.

. . .

I have been trying to setup a multi node Kubernetes cluster to try out Ballerina integrations at scale in containers.

There are many third party tools to setup a Kubernetes cluster. I've looked at two options, Minikube and Kubeadm. Minikube doesn't support multi node setup, so I picked Kubeadm. Kubeadm is still in beta, do not use it in production.

I followed Kubeadm guide and setup a three node cluster (single master and two nodes) manually first and then automated in Vagrant because I can recreate the cluster easily later with no effort.

So here is how I did it. Try it, create issues in my repo and send pull requests.

**Welcome to VirtualBox!**

The left part of this window lists all virtual machines and virtual machine groups on your computer. The list is empty now because you haven't created any virtual machines yet.

In order to create a new virtual machine, press the **New** button in the main tool bar located at the top of the window.

**1. vagrant@k8s-head: ~ (ssh)**

```

vagrant@k8s-head:~$ kubectl get pods
NAME                                READY     STATUS    RESTARTS   AGE
kube-system/calico-node-fflq9       2/2      Running   0           30m
kube-system/calico-node-gvfgd       2/2      Running   0           37m
kube-system/coredns-78fcd6894-75sdc 1/1      Running   0           37m
kube-system/coredns-78fcd6894-vjlsm 1/1      Running   0           37m
kube-system/etcd-k8s-head           1/1      Running   0           36m
kube-system/kube-apiserver-k8s-head 1/1      Running   0           36m
kube-system/kube-controller-manager-k8s-head 1/1      Running   0           36m
kube-system/kube-proxy-lk4w9        1/1      Running   0           37m
kube-system/kube-proxy-mdr7b        1/1      Running   0           24m
kube-system/kube-proxy-vc4wk        1/1      Running   0           30m
kube-system/kube-scheduler-k8s-head 1/1      Running   0           36m
vagrant@k8s-head:~$

```

**vagrant@k8s-node-1: ~ (ssh)**

PORTS	NAMES	AGE	STATUS
d9760cff9446	quay.io/calico/cni@sha256:ed172c28bc193bb09bce6be6ed7dc6bfc85118d55e61d263cee8bbb0fd464a9d	30 minutes ago	Up 29 minut
es	k8s_install-cni_calico-node-fflq9_kube-system_7eb64021-b127-11e8-aba6-021d76b33cfe_0		
0ea666a3c575	k8s.gcr.io/kube-proxy-amd64@sha256:6a8d6e8d1674cb26167d85bebbb953e93993b81bbbf7e00c2985e61e0c7c2062	30 minutes ago	Up 30 minut
es	k8s_kube-proxy_kube-proxy-vc4wk_kube-system_7eb6ed02-b127-11e8-aba6-021d76b33cfe_0		
8e31cae49343	quay.io/calico/node@sha256:a35541153f7695b38afada46843c64a2c546548cd8c171f402621736c6cf3f0b	30 minutes ago	Up 30 minut
es	k8s_calico-node_calico-node-fflq9_kube-system_7eb64021-b127-11e8-aba6-021d76b33cfe_0		
e0e5d6725441	k8s.gcr.io/pause:3.1	30 minutes ago	Up 30 minut
es	k8s_POD_calico-node-fflq9_kube-system_7eb64021-b127-11e8-aba6-021d76b33cfe_0		
dc9c9dcfe029	k8s.gcr.io/pause:3.1	30 minutes ago	Up 30 minut
es	k8s_POD_kube-proxy-vc4wk_kube-system_7eb6ed02-b127-11e8-aba6-021d76b33cfe_0		

**vagrant@k8s-node-1:~\$**

**vagrant@k8s-node-2: ~ (ssh)**

PORTS	NAMES	AGE	STATUS
a1954f583a3d	quay.io/calico/cni@sha256:ed172c28bc193bb09bce6be6ed7dc6bfc85118d55e61d263cee8bbb0fd464a9d	23 minutes ago	Up 23 minut
es	k8s_install-cni_calico-node-5hrp1_kube-system_6355c815-b128-11e8-aba6-021d76b33cfe_0		
602d0e88995b	k8s.gcr.io/kube-proxy-amd64@sha256:6a8d6e8d1674cb26167d85bebbb953e93993b81bbbf7e00c2985e61e0c7c2062	23 minutes ago	Up 23 minut
es	k8s_kube-proxy_kube-proxy-mdr7b_kube-system_635435ea-b128-11e8-aba6-021d76b33cfe_0		
4522f3acea28	quay.io/calico/node@sha256:a35541153f7695b38afada46843c64a2c546548cd8c171f402621736c6cf3f0b	23 minutes ago	Up 23 minut
es	k8s_calico-node_calico-node-5hrp1_kube-system_6355c815-b128-11e8-aba6-021d76b33cfe_0		
759195231c85	k8s.gcr.io/pause:3.1	24 minutes ago	Up 24 minut
es	k8s_POD_calico-node-5hrp1_kube-system_6355c815-b128-11e8-aba6-021d76b33cfe_0		
c972e889841b	k8s.gcr.io/pause:3.1	24 minutes ago	Up 24 minut

**vagrant@k8s-node-2:~\$**



*Note: Step-1 to Step-8 should be done in all the Kubernetes nodes (master + nodes). Step-9 to Step-12 should be done in Kubernetes master. Step-13 should be done in Kubernetes nodes.*

### Step-1: Pick a base box for Kubernetes nodes

```
config.vm.box = "ubuntu/xenial64"  
config.vm.box_version = "20180831.0.0"
```

### Step-2: Setup a network

```
config.vm.network :private_network, ip: "192.168.205.10"
```

### Step-3: Setup memory and CPUs

```
v.customize ["modifyvm", :id, "--memory", "2048"]  
v.customize ["modifyvm", :id, "--cpus", "2"]
```

## Step-4: Install Docker

I'm not going to use vagrant docker provisioning because it always installs the latest version, but Kubeadm requires 17.03 or older. So I'm going to use shell provisioning to install docker.

```
apt-get update
apt-get install -y apt-transport-https ca-certificates curl
software-properties-common
curl -fsSL https://download.docker.com/linux/ubuntu/gpg |
apt-key add -
add-apt-repository "deb
https://download.docker.com/linux/$(. /etc/os-release; echo
"$ID") $(lsb_release -cs) stable"
apt-get update && apt-get install -y docker-ce=$(apt-cache
madison docker-ce | grep 17.03 | head -1 | awk '{print $3}')
```

## Step-5: Allow regular user to run docker commands

I'm adding the regular user (vagrant) to the docker user group so that I can run docker commands as vagrant user without sudo privileges.

```
usermod -aG docker vagrant
```

## Step-6: Install Kubeadm, Kubelet and Kubctl

I'm using shell provisioning to install Kubeadm, Kubelet and Kubectl from the repositories.

```
apt-get install -y apt-transport-https curl
curl -s https://packages.cloud.google.com/apt/doc/apt-
```

```
key.gpg | apt-key add -  
cat <<EOF >/etc/apt/sources.list.d/kubernetes.list  
deb http://apt.kubernetes.io/ kubernetes-xenial main  
EOF  
apt-get update  
apt-get install -y kubelet kubeadm kubectl  
apt-mark hold kubelet kubeadm kubectl
```

## Step-7: Kubelet requires swap to be disabled

Running following bash command will disable the swap in the current login session.

```
swapoff -a
```

## Step-8: Disable swap permanently

Following change will keep swap off after reboot too

```
sudo sed -i ' / swap / s/^(.*)$/#\1/g' /etc/fstab
```

## Step-9: Initialize Kubernetes master

Following Kubeadm command will install and configure a Kubernetes master.

```
# ip of this box  
IP_ADDR=$(ifconfig enp0s8 | grep Mask | awk '{print $2}' | cut  
-f2 -d:)
```

```
# install k8s master
HOST_NAME=$(hostname -s)
kubeadm init --apiserver-advertise-address=$IP_ADDR --
apiserver-cert-extra-sans=$IP_ADDR --node-name $HOST_NAME -
-pod-network-cidr=192.168.0.0/16
```

## Step-10: Allow regular user to run Kubectl commands

I'm going to copy the Kubectl credentials to vagrant user's home directory to be able to run Kubectl commands without sudo privileges.

```
#copying credentials to regular user - vagrant
sudo --user=vagrant mkdir -p /home/vagrant/.kube
cp -i /etc/kubernetes/admin.conf /home/vagrant/.kube/config
chown $(id -u vagrant):$(id -g vagrant)
/home/vagrant/.kube/config
```

## Step-11: Install a Pod Network

We need to install a Pod Network to enable pod-to-pod communication in our Kubernetes cluster. There are many third party Pod Network projects such as Calico, Flannel and Weave. I've tried to use Weave, but looks like there are some bugs that prevent a Weave pod to be scheduled on Kubernetes nodes. Then I've tried Calico and it works fine so far.

```
# install Calico pod network addon
export KUBECONFIG=/etc/kubernetes/admin.conf
kubectl apply -f
https://docs.projectcalico.org/v3.1/getting-
started/kubernetes/installation/hosted/rbac-kdd.yaml
kubectl apply -f
```

```
https://docs.projectcalico.org/v3.1/getting-  
started/kubernetes/installation/hosted/kubernetes-  
datastore/calico-networking/1.7/calico.yaml
```

## Step-12: Save the Kubeadm join commands

Our Kubernetes master should be up and running by now. We can join Kubernetes nodes to master using Kubeadm join command. You can run `kubeadm token create --print-join-command` in Kubernetes master to get the join command that should be executed in Kubernetes nodes.


I'm going to store the output of this command in a file in Kubernetes master. Later I'm going to download this file from inside Kubernetes nodes, execute and join the cluster.

```
kubeadm token create --print-join-command >>  
/etc/kubeadm_join_cmd.sh  
chmod +x /etc/kubeadm_join_cmd.sh
```

## Step-13: Execute Kubeadm join command in Kubernetes nodes

What I'm doing here is to scp the above `kubeadm_join_cmd.sh` from master to nodes and execute it. I'm using `sshpass` to run prompt-less scp command in bash.

```
apt-get install -y sshpass  
sshpass -p "vagrant" scp -o StrictHostKeyChecking=no  
vagrant@192.168.205.10:/etc/kubeadm_join_cmd.sh .  
sh ./kubeadm_join_cmd.sh
```



That's all, your cluster should be up and running now.

### Putting it all together

Here is my full Vagrantfile gist. But be sure check [my repo](#) to get the up-to-date Vagrant file. Run `vagrant up` to bring up a 3 node Kubernetes cluster.



```
1  # -*- mode: ruby -*-
2  # vi: set ft=ruby :
3
4  servers = [
5    {
6      :name => "k8s-head",
7      :type => "master",
8      :box => "ubuntu/xenial64",
9      :box_version => "20180831.0.0",
10     :eth1 => "192.168.205.10",
11     :mem => "2048",
12     :cpu => "2"
13   },
14   {
15     :name => "k8s-node-1",
16     :type => "node",
17     :box => "ubuntu/xenial64",
18     :box_version => "20180831.0.0",
19     :eth1 => "192.168.205.11",
20     :mem => "2048",
21     :cpu => "2"
22   },
23   {
24     :name => "k8s-node-2",
25     :type => "node",
26     :box => "ubuntu/xenial64",
27     :box_version => "20180831.0.0",
28     :eth1 => "192.168.205.12",
29     :mem => "2048",
30     :cpu => "2"
31   }
32 ]
33
34 # This script to install k8s using kubeadm will get e
```

```

35 $configureBox = <<-SCRIPT
36
37     # install docker v17.03
38     # reason for not using docker provision is that i
39     apt-get update
40     apt-get install -y apt-transport-https ca-certifi
41     curl -fsSL https://download.docker.com/linux/ubun
42     add-apt-repository "deb https://download.docker.c
43     apt-get update && apt-get install -y docker-ce=$(
44
45     # run docker commands as vagrant user (sudo not r
46     usermod -aG docker vagrant
47
48     # install kubeadm
49     apt-get install -y apt-transport-https curl
50     curl -s https://packages.cloud.google.com/apt/doc
51     cat <<EOF >/etc/apt/sources.list.d/kubernetes.lis
52     deb http://apt.kubernetes.io/ kubernetes-xenial m
53 EOF
54     apt-get update
55     apt-get install -y kubelet kubeadm kubectl
56     apt-mark hold kubelet kubeadm kubectl
57
58     # kubelet requires swap off
59     swapoff -a
60
61     # keep swap off after reboot
62     sudo sed -i '/ swap / s/^(.*)$/#\1/g' /etc/fsta
63
64     # ip of this box
65     IP_ADDR=`ifconfig enp0s8 | grep Mask | awk '{prin
66     # set node-ip
67     sudo sed -i "/^[^#]*KUBELET_EXTRA_ARGS=/c\KUBELET
68     sudo systemctl restart kubelet

```

```
69  SCRIPT
70
71  $configureMaster = <<--SCRIPT
72      echo "This is master"
73      # ip of this box
74      IP_ADDR=`ifconfig enp0s8 | grep Mask | awk '{prin
75
76      # install k8s master
77      HOST_NAME=$(hostname -s)
78      kubeadm init --apiserver-advertise-address=$IP_AD
79
80      #copying credentials to regular user - vagrant
81      sudo --user=vagrant mkdir -p /home/vagrant/.kube
82      cp -i /etc/kubernetes/admin.conf /home/vagrant/.k
83      chown $(id -u vagrant):$(id -g vagrant) /home/vag
84
85      # install Calico pod network addon
86      export KUBECONFIG=/etc/kubernetes/admin.conf
87      kubectl apply -f https://docs.projectcalico.org/v
88      kubectl apply -f https://docs.projectcalico.org/v
89
```



